

## Lecture 19: Kernel PCA and MLE

Nov 7th 2019

Lecturer: Steven Wu

Scribe: Steven Wu

## Principal Component Analysis

*Principal component analysis* aims to solve the following optimization problem: given as input a data matrix  $X \in \mathbb{R}^{n \times d}$ , find an encoder  $E$  and decoder  $D$  to minimize the following reconstruction error:

$$\min_{D \in \mathbb{R}^{k \times d}, E \in \mathbb{R}^{d \times k}} \|X - XED\|_F^2 \quad (1)$$

where  $\|\cdot\|_F$  denotes the Frobenius norm: for any matrix  $A$ ,

$$\|A\|_F = \sqrt{\sum_{(i,j)} A_{ij}^2} = \sqrt{\text{tr}(A^T A)}$$

The PCA method solves the problem with the following procedure: compute  $X = USV^T$ , then return encoder  $E = V_k$ , decoder  $D = V_k^T$ , encoded data  $XV_k = U_k S_k \in \mathbb{R}^{n \times k}$ , and decoded data  $XV_k V_k^T$ . Note that  $V_k V_k^T \in \mathbb{R}^{d \times d}$  performs orthogonal projection onto subspace spanned by  $V_k$ .

Last lecture, we showed that the optimization problem can be re-written as

$$\min_{D \in \mathbb{R}^{k \times d}, E \in \mathbb{R}^{d \times k}} \|X - XED\|_F^2 = \min_{D \in \mathbb{R}^{d \times k}, D^T D = I} \|X - XDD^T\|_F^2$$

We also showed that this new objective can be further decomposed

$$\|X - XDD^T\|_F^2 = \|X\|_F^2 - \|XD\|_F^2$$

This means,

$$\min_{D \in \mathbb{R}^{d \times k}, D^T D = I} \|X - XDD^T\|_F^2 \Leftrightarrow \max_{D \in \mathbb{R}^{d \times k}, D^T D = I} \|XD\|_F^2$$

Finally, the objective value of the maximization problem is singular values squared.

$$\max_{D \in \mathbb{R}^{d \times k}, D^T D = I} \|XD\|_F^2 = \|XV_k\|_F^2 = \sum_{j=1}^k s_j^2$$

where  $s_1, \dots, s_k$  are the top singular values of  $X$ .

**Centered PCA.** Typically, before running PCA, we replace each  $x_i$  with  $x'_i = x_i - \bar{x}$ , where  $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$ . The objective then becomes

$$\|X'D\|_F^2 = \text{tr}((X'D)^\top(X'D)) = \sum_{i=1}^k (X'De_i)^\top(X'De_i)$$

Note that  $\frac{1}{n}(X'De_i)^\top(X'De_i)$  corresponds to the variance on the  $i$ -th coordinate after the projection. Therefore, PCA is maximizing the resulting per-coordinate variances.

**Power method.** How to compute the SVD of  $X \in \mathbb{R}^{n \times d}$ ? We can use the power method to first compute  $v_1$ ,  $u_1$  and  $s_1$ . The idea is to compute the top eigenvector and eigenvalue of the matrix  $X^\top X$ :

- Start with a random vector  $y_0 \sim \mathcal{N}(0, I_d)$
- For  $t = 1, \dots, T$ :  
 $y_t \leftarrow X^\top X y_{t-1}$
- $v_1 \leftarrow y_T / \|y_T\|_2$ : the first column of  $V$  and also the top eigenvector of  $X^\top X$
- $s_1 \leftarrow \|Xv_1\|_2$  top singular value
- $u_1 \leftarrow Xv_1 / s_1$

To compute the remainder of triplets  $(u_i, s_i, v_i)$ , repeat the same for the residual matrix  $X - s_1 u_1 v_1^\top$ . Note that we can also apply the power method to the matrix  $XX^\top$  for computing its top eigenvector, which is  $u_1$ . This will be useful for the next kernel PCA method.

## Kernel PCA

We can find the “high variance” directions in a richer feature space by first apply some feature mapping  $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^m$  and then runs PCA over the transformed data. Let  $\Phi \in \mathbb{R}^{n \times m}$  such that each row of  $\Phi$  is given by  $\phi(x_i)$ . Let  $k(\cdot, \cdot)$  be the kernel such that  $k(x, y) = \phi(x)^\top \phi(y)$ . Kernel PCA then does the following:

- Compute the Gram matrix  $G = \Phi\Phi^\top$  and the centered Gram matrix

$$\begin{aligned} \overline{G} &= (\Phi - E\Phi)(\Phi - E\Phi)^\top \\ &= \Phi\Phi^\top - E\Phi\Phi^\top - \Phi\Phi^\top E + E\Phi\Phi^\top E \\ &= G - EG - GE + EGE \end{aligned}$$

where  $E \in \mathbb{R}^{n \times n}$  is the matrix with all entries of  $1/n$ .

- Find the top  $k$  eigenvectors of  $\overline{G}$  with normalization: call it  $A \in \mathbb{R}^{n \times k}$




Figure 1: Denoising application of kernel PCA on the digits data set. Image from Haipeng Luo's lecture slide. Another application [here](#).

- Construct the encoded dataset

$$(\Phi - E\Phi)(\Phi - E\Phi)^T A = \bar{G}A$$

## Maximum Likelihood Estimation

Now we will switch over to a different unsupervised learning problem that aims to model the underlying probability distribution  $P$  from which the observed examples are drawn. As we discussed before in the lecture on logistic regression, there is a general principle called *maximum likelihood estimation* (MLE):

- Pick a set of probability models for the data:  $\mathcal{P} := \{p_\theta : \theta \in \Theta\}$ .

- Given samples  $x_1, \dots, x_n$ , pick the model that maximized the likelihood

$$\max_{\theta \in \Theta} L(\theta) = \max_{\theta \in \Theta} \ln \prod_{i=1}^n p_{\theta}(x_i) = \max_{\theta \in \Theta} \sum_{i=1}^n \ln p_{\theta}(x_i)$$

**Example 0.1** (Coin flips). *Heads:  $x_i = 1$  and tails  $x_i = 0$ . The Bernoulli distribution has the parameter—the bias or the probability of heads  $\theta \in [0, 1]$ . We can write*

$$p_{\theta}(x_i) = \theta^{x_i} (1 - \theta)^{1-x_i}$$

Let  $H = \sum_i x_i$  and  $T = \sum_i (1 - x_i)$  be the number of heads and tails.

$$L(\theta) = \sum_{i=1}^n (x_i \ln \theta + (1 - x_i) \ln(1 - \theta)) = H \ln \theta + T \ln(1 - \theta).$$

By using the first-order condition, we derive that the solution is

$$\theta = \frac{H}{T + H} = \frac{H}{n}.$$

## Gaussian Mixture Model

Gaussian mixture model (GMM) is the following generative model:

- Draw a latent class  $Y$  such that  $\Pr[Y = j] = \pi_j$
- Then draw  $X$  conditioned on  $Y$ :  $X \mid Y = j \sim \mathcal{N}(\mu_j, \Sigma_j)$ .

The parameter  $\theta = ((\pi_1, \mu_1, \Sigma_1), \dots, (\pi_k, \mu_k, \Sigma_k))$  and the probability density at each point  $x$  is

$$p_{\theta}(x) = \sum_{j=1}^k p_{\mu_j, \Sigma_j}(x) \pi_j$$

where  $p_{\mu_j, \Sigma_j}$  denotes the multivariate Gaussian density function:

$$p_{\mu_j, \Sigma_j}(x) = \frac{1}{\sqrt{(2\pi)^d \det(\Sigma_j)}} \exp \left( -\frac{1}{2} (x - \mu_j)^T \Sigma_j^{-1} (x - \mu_j) \right)$$

The MLE problem becomes minimization of

$$L(\theta) = \sum_{i=1}^n \ln \left( \sum_{j=1}^k p_{\mu_j, \Sigma_j}(x_i) \pi_j \right)$$




Figure 2: Statistician Karl Pearson wanted to understand the distribution of the ratio between forehead breadth and body length for crabs. He fit a mixture of two Gaussians. Figure due to Peter Macdonald.