

## Lecture 21: Variational Autoencoders

Nov 14th 2019

Lecturer: Steven Wu

Scribe: Steven Wu

Now we will study how to leverage generative models to *sample* from a distribution. We will leverage neural networks in the following way:

- First sample a latent variable  $z$  from distribution  $\mu$ , that is easy to sample from. For example,  $\mu$  can be the uniform distribution over  $[0, 1]$  or the Gaussian distribution.
- Then pass the latent variable through a neural network  $g$  and output  $g(z)$ .

In this lecture, we will cover one of the most popular generative network method—*variational autoencoder* (VAE).

**Autoencoder** Let us first talk about what an autoencoder is. Well, in fact, you have already seen an autoencoder at this point. A special case is just the PCA (and also kernel PCA), which gives the optimal linear encoding/decoding: Given  $X = USV^T$  and  $k \leq r$ ,

$$\min_{E \in \mathbb{R}^{d \times k}, D \in \mathbb{R}^{k \times d}} \|X - XED\|_F^2 = \|X - XV_k V_k^T\|_F^2$$

But we can also have encoders and decoders that are not linear mappings. Let encoders  $\mathcal{E}$  and decoders  $\mathcal{D}$  denote families of deep networks from  $\mathbb{R}^d$  to  $\mathbb{R}^k$  and from  $\mathbb{R}^k$  to  $\mathbb{R}^d$

$$\min_{f \in \mathcal{E}, g \in \mathcal{D}} \sum_{i=1}^n \|x_i - g(f(x_i))\|_2^2$$

This is called an autoencoder, which deterministically map each example  $x_i$  to a latent code  $z_i$ , back to some approximation of  $x_i$ . We say that  $\mathbb{R}^k$  is the latent space, and  $f(x) \in \mathbb{R}^k$  is latent representation of  $x$ .

## Variational Autoencoder (VAE)

We will now leverage the idea of autoencoder to build generative models. Intuitively, we should take the decoder  $g$  from an autoencoder as our generative network, which is a mapping from a low-dimensional latent space  $\mathbb{R}^k$  to the example space  $\mathbb{R}^d$ .

In particular, suppose we have a sample  $x_1, \dots, x_n$  drawn from some distribution  $P$ . We want to find  $g$  so that  $g(z_i) \approx x_i$  for each  $i$ , where each  $z_i$  is drawn from a Gaussian distribution. VAE construct a distribution for each  $z_i$  based on each  $x_i$ . The method runs over iterations, and in each iteration does the following:

1. Encode each example into Gaussian mean-variance parameters  $(\mu_i, \Sigma_i) \leftarrow f(x_i)$ .
2. Sample latent variable from Gaussian:  $z_i \sim N(\mu_i, \Sigma_i)$ .
3. Decode  $\hat{x}_i = g(z_i)$ .
4. Taking a gradient descent step (or any other optimization method) to further minimize the VAE objective

$$\sum_{i=1}^n \ell(x_i, \hat{x}_i) + \lambda \text{KL}(\mathcal{N}(\mu_i, \sigma_i^2 I), \mathcal{N}(0, I))$$

where  $\ell(x_i, \hat{x}_i)$  is “reconstruction error”. For example,  $\ell(x_i, \hat{x}_i) = \|x_i - \hat{x}_i\|_2^2$ . We will go into the details of the gradient update step in a bit.

In the VAE objective, KL denotes KL divergence: for any two distributions  $p$  and  $q$ ,

$$\text{KL}(p||q) = \int p(z) \ln \frac{p(z)}{q(z)} dz$$

KL divergence is a dissimilarity measure between distributions, with two important properties:

- $\text{KL}(p||q) \geq 0$  for any  $p, q$ .
- $\text{KL}(p||q) = \text{KL}(q||p)$  if and only if  $p = q$ .

KL divergence encourages the individual distributions  $\mathcal{N}(\mu_i, \Sigma_i)$  to be close to the distribution  $\mathcal{N}(0, I)$ . This is useful because  $\mathcal{N}(0, I)$  is the “source” distribution for the generative models—that is, we output  $g(z)$  with  $z \sim \mathcal{N}(0, I)$ . The smaller the KL divergence is, the closer this sampling has to approximate the training distribution.

## Derivation from Variational Inference

VAE is based on ideas from *variational inference* (VI), which is a popular method to perform approximate inference in probabilistic models. We won’t get into the details of VI here, but we will discuss the relevant ideas that lead to VAE. Let  $\mathcal{P} = \{p_\theta \mid \theta \in \Theta\}$  be a family of probability distributions over observed and latent variables  $x$  and  $z$ . Given a set of observed variables  $S = \{x_1, \dots, x_n\}$ , we would like to find a distribution in  $\mathcal{P}$  to minimize:

$$\min_{p \in \mathcal{P}} \text{KL}(\hat{p}_S || p) = \min_{p \in \mathcal{P}} \sum_{x \in S} \hat{p}_S(x) \ln \frac{\hat{p}_S(x)}{p(x)}$$

where  $\hat{p}_S$  denotes the empirical distribution over the data set. Note that  $\sum_{x \in S} \hat{p}_S(x) \ln p_s(x)$  does not depend on the choice of  $p$ . Thus, the minimization is equivalent to the following maximization problem:

$$\max_{p \in \mathcal{P}} \sum_{x \in S} \hat{p}_S(x) \ln p(x) \Leftrightarrow \max_{p \in \mathcal{P}} \sum_{x_i \in S} \ln p(x_i) \Leftrightarrow \max_{p \in \mathcal{P}} \sum_{x_i \in S} \ln \int p(x_i, z) dz$$

latent  $z$   $\longrightarrow$  observed  $x$

Figure 1: Graphical model with latent variable

Thus, minimizing the KL divergence objective is the same as maximizing log-likelihood. The problem above is typically intractable for generative models with high-dimensional  $z$ , since it involves computing an integral over all  $z$ 's.

To circumvent the intractability, the VI method aims to optimize a tractable lower bound of the log-likelihood. To do that, we introduce a family of approximate distributions  $\mathcal{Q} = \{q_\gamma \mid \gamma \in \Gamma\}$ . (Each distribution  $q$  is parameterized by  $\gamma$ .) Observe that for any fixed  $x$ ,

$$\begin{aligned} \ln p(x) &= \int q(z|x) \ln p(x) dz \\ &= \int q(z|x) \ln \frac{p(x)q(z|x)p(z|x)}{p(z|x)q(z|x)} dz \\ &= \int q(z|x) \ln \frac{q(z|x)}{p(z|x)} dz + \int q(z|x) \ln \frac{p(x, z)}{q(z|x)} dz \\ &= \underbrace{\text{KL}(q(z|x) \parallel p(z|x))}_{\geq 0} + \underbrace{\int q(z|x) \ln \frac{p(x, z)}{q(z|x)} dz}_{\text{ELBO}} \end{aligned}$$

As indicated above, the KL term is always non-negative, and so the second term is a lower bound for  $\ln p(x)$ . The second term is hence called the *evidence lower bound* (ELBO). For any two distributions  $p_\theta \in \mathcal{P}$  and  $q_\gamma \in \mathcal{Q}$ , let us write

$$\text{ELBO}(x; \theta, \gamma) = \int q(z|x) \ln \frac{p_\theta(x, z)}{q_\gamma(z|x)} dz$$

The VI method then uses gradient-based method to optimize the objective

$$\max_{\theta} \sum_{x_i \in S} \max_{\gamma_i} \mathbf{E}_{q_{\gamma_i}(z|x_i)} \left[ \log \frac{p_\theta(x_i, z)}{q_{\gamma_i}(z|x_i)} \right]. \quad (1)$$

In each iteration, we do two-step update:

1. First, for each example  $i$ : update  $\gamma_i$

$$\gamma_i \leftarrow \gamma_i + \eta_\gamma \tilde{\nabla}_\gamma \text{ELBO}(x_i; \theta, \gamma^{(i)}), \quad (2)$$

2. Update  $\theta$

$$\theta \leftarrow \theta + \eta_\theta \tilde{\nabla}_\theta \sum_i \text{ELBO}(x^{(i)}; \theta, \gamma^{(i)}), \quad (3)$$

where  $\tilde{\nabla}$  denote unbiased estimate for the gradients and  $\eta_\gamma$  and  $\eta_\theta$  are the learning rates.

**Reparameterization trick.** To estimate the gradient  $\nabla_{\gamma} \text{ELBO}(x; \theta, \gamma) = \nabla_{\gamma} \mathbf{E}_{q_{\gamma}(z|x)} \left[ \log \frac{p_{\theta}(x, z)}{q_{\gamma}(z|x)} \right]$ , we will leverage a *reparameterization trick*. Let us introduce a fixed, auxiliary distribution  $\nu(\epsilon)$  and a differentiable function  $T(\epsilon; \gamma)$  such that sampling from  $q_{\gamma}(z|x)$  is identical to

$$\epsilon \sim \nu \quad z \sim T(\epsilon; \gamma)$$

Then the gradient computation can be rewritten as:

$$\nabla_{\gamma} \mathbf{E}_{q_{\gamma}(z|x)} \left[ \log \frac{p_{\theta}(x, z)}{q_{\gamma}(z|x)} \right] = \mathbf{E}_{\nu} \left[ \nabla_{\gamma} \log \frac{p_{\theta}(x, T(\epsilon; \gamma))}{q_{\gamma}(T(\epsilon; \gamma))} \right] \quad (4)$$

We can then approximate the right hand side of (4) by drawing  $\epsilon_1, \dots, \epsilon_m$  from  $\nu$ , and then compute the average gradient:

$$\frac{1}{m} \sum_{i=1}^m \left[ \nabla_{\gamma} \log \frac{p_{\theta}(x, T(\epsilon_i; \gamma))}{q_{\gamma}(T(\epsilon_i; \gamma))} \right]$$

This is also called Monte Carlo sampling. Note that the gradient  $\nabla_{\theta} \text{ELBO}(x; \theta, \gamma)$  can be estimated with Monte Carlo sampling, but without the reparametrization trick: draw  $z_1, \dots, z_m$  i.i.d. from  $p(z|x)$ , and then compute the average gradient

$$\frac{1}{m} \sum_{i=1}^m \left[ \nabla_{\theta} \log \frac{p_{\theta}(x, z_i)}{q_{\gamma}(z_i|x)} \right]$$

where  $\Sigma^{1/2}$  is the Cholesky decomposition of  $\Sigma$ .

**Instantiation via neural nets.** Now we will obtain VAE from this framework of VI by instantiating the distributions  $p$  and  $q$  through neural networks and Gaussian distributions. First, we will have the latent distribution as

$$p_{\theta}(z) = \mathcal{N}(0, I)$$

Note that this “prior” distribution doesn’t depend on  $\theta$ . The conditional distribution  $p_{\theta}(x|z)$  corresponds to the decoder. A typical choice is a Gaussian distribution

$$p_{\theta}(x|z) = \mathcal{N}(\mu_{\theta}(z), \Sigma_{\theta}(z))$$

where the mean and covariance parameters  $\mu_{\theta}(z), \Sigma_{\theta}(z)$  are given by a neural network. If  $\Sigma_{\theta}(z) = \sigma^2 I$ , then ELBO becomes the VAE objective with squared error as the reconstruction error, that is

$$\ell(x_i, \hat{x}_i) = \|x_i - \hat{x}_i\|_2^2$$

For the approximate distribution  $q$ , we will have

$$q_{\gamma}(z|x_i) = \mathcal{N}(\mu(x_i), \Sigma(x_i)),$$

where the parameter  $\gamma_i = (\mu(x_i), \Sigma(x_i))$  are mean and covariance parameters given by the encoder neural network. To apply the reparameterization trick, we will have  $\nu = \mathcal{N}(0, I)$  and  $T(\epsilon; \gamma) = \mu + \Sigma^{1/2} \epsilon$ , where  $\Sigma^{1/2}$  is the Cholesky decomposition of  $\Sigma$ . For  $\Sigma = \sigma^2 I$ , we will simply have  $\Sigma^{1/2} = \sigma I$ .